

# Predicting Customer Lifetime Value - Beginner

---

[KeytoDataScience.com](https://keytoDataScience.com)

For any Data Science project, we will follow the approach of Data Science Project Lifecycle:

---

## **Data Science Project Life Cycle** <sup>1</sup>

---

1. **Understand Problem/Objective**
  2. **Data Collection**
  3. **Data Preparation**
    - 3.1 Data Preprocessing
    - 3.2 EDA<sup>3</sup>
    - 3.3 Train/Validation/Test Split
    - 3.4 Feature Engineering
    - 3.5 Feature Selection
  4. **Modeling:** Regression <sup>6</sup>
  5. **Evaluation:** Regression
    - RMSE, RSE, MAE, RAE, Coefficient of Determination (R2)
  6. **Model Deployment**
    - Model Deployment in pipeline or tool
- 

In this case study, we will use past purchase history of customers to build a model that can predict the **Customer Lifetime Value (CLV or CLTV)** for new customers.

## Table of Contents

- [1. Understand Objective](#)
- [2. Data Collection](#)
- [3. Data Preparation](#)
  - [3.1 Data Preprocessing](#)
  - [3.2 Exploratory Data Analysis \(EDA\)](#)
  - [3.3 Train and Test Split](#)
- [4. Modeling](#)
  - [4.1 Linear Regression](#) <sup>6</sup>
  - [4.2 DecisionTreeRegressor](#) <sup>7</sup>
  - [4.3 RandomForestRegressor](#) <sup>8</sup>

- 5. Evaluation <sup>9</sup>
  - 5.1 Select Final Model
  - 5.2 Save Model to Disk
  - 5.3 Interpret the Output
  - 5.4 Linear Regression with StandardScaler (Optional)
- 6. Model Deployment
  - 6.1 Import Libraries
  - 6.2 Load Model from Disk
  - 6.3 Real Time Prediction

# 1. Understand Objective

## Customer Lifetime Value(CLTV)<sup>2</sup>

"Customer Lifetime Value is a monetary value that represents the amount of revenue or profit a customer will give the company over the period of the relationship" (Source). CLTV demonstrates the implications of acquiring long-term customers compare to short-term customers. Customer lifetime value (CLV) can help you to answers the most important questions about sales to every company:

1. How to Identify the most profitable customers?
2. How can a company offer the best product and make the most money?
3. How to segment profitable customers?
4. How much budget need to spend to acquire customers?

## Business Terms

- **Average Order Value(AOV):** The Average Order value is the ratio of your total revenue and the total number of orders. AOV represents the mean amount of revenue that the customer spends on an order.
  - $\text{Average Order Value} = \text{Total Revenue} / \text{Total Number of Orders}$
- **Purchase Frequency:** Purchase Frequency is the ratio of the total number of orders and the total number of customer. It represents the average number of orders placed by each customer.
  - $\text{Purchase Frequency} = \text{Total Number of Orders} / \text{Total Number of Customers}$
- **Churn Rate:** Percentage of customers who have not ordered again.
- **Customer Lifetime:** Customer Lifetime is the period of time that the customer has been continuously ordering.
  - $\text{Customer lifetime} = 1 / \text{Churn Rate}$
- **Repeat Rate:** Repeat rate can be defined as the ratio of the number of customers with more than one order to the number of unique customers. Example: If you have 10 customers in a month out of who 4 come back, your repeat rate is 40%.

- Repeat Rate = 1 - Churn Rate

[↑ back to top](#)

## 2. Data Collection

### Import Libraries

```
In [1]: from pandas import Series, DataFrame
import pandas as pd
import numpy as np
import os

import matplotlib.pyplot as plt
plt.style.use('ggplot')
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler

import sklearn.metrics
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

import joblib
```

### Download Data<sup>2</sup>

#### Load Data

```
In [2]: df = pd.read_csv("history.csv")
```

We will load the data file for this example and checkout summary statistics and columns for that file.

#### Check out the Data

```
In [3]: df.shape
```

```
Out[3]: (100, 8)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CUST_ID     100 non-null   int64
1   MONTH_1     100 non-null   int64
2   MONTH_2     100 non-null   int64
3   MONTH_3     100 non-null   int64
4   MONTH_4     100 non-null   int64
5   MONTH_5     100 non-null   int64
```

```

6  MONTH_6  100 non-null    int64
7  CLV      100 non-null    int64
dtypes: int64(8)
memory usage: 6.4 KB

```

The dataset consists of the customer ID, the amount the customer spent on your website for the first months of his relationship with your business and his ultimate life time value ( say 3 years worth)

In [5]: `df.describe()`

```

Out[5]:
      CUST_ID  MONTH_1  MONTH_2  MONTH_3  MONTH_4  MONTH_5  MONTH_6  CLV
count  100.000000    100.00000    100.000000    100.000000    100.000000    100.000000    100.000000    100.000000
mean    1050.500000    113.25000    115.750000    106.250000    106.750000    106.250000    108.500000    9421.190000
std      29.011492     55.32902     64.221739     63.601406     62.649317     59.816111     66.021499    2664.44317
min    1001.000000     25.00000     0.000000     0.000000     0.000000     0.000000     0.000000    4125.000000
25%    1025.750000     75.00000    75.000000    50.000000    50.000000    50.000000    50.000000    7816.000000
50%    1050.500000    100.00000   125.000000   100.000000   100.000000   100.000000   100.000000    9344.000000
75%    1075.250000   150.00000   175.000000   175.000000   150.000000   156.250000   175.000000   10719.250000
max    1100.000000   200.00000   200.000000   200.000000   200.000000   200.000000   200.000000   17100.000000

```

In [6]: `df.head()`

```

Out[6]:
      CUST_ID  MONTH_1  MONTH_2  MONTH_3  MONTH_4  MONTH_5  MONTH_6  CLV
0      1001         150         75         200         100         175         75  13125
1      1002          25         50         150         200         175         200  9375
2      1003          75        150          0          25          75          25  5156
3      1004         200        200          25         100          75         150  11756
4      1005         200        200        125          75         175         200  15525

```

[↑ back to top](#)

## 3. Data Preparation

### 3.1 Data Preprocessing

```

In [7]:
# drop CUST_ID
df=df.drop("CUST_ID",axis=1)

# or
# df.drop("CUST_ID",axis=1,inplace=True)

```

## Check Null Values

```
In [8]: df.isnull().sum()
```

```
Out[8]: MONTH_1    0
        MONTH_2    0
        MONTH_3    0
        MONTH_4    0
        MONTH_5    0
        MONTH_6    0
        CLV        0
        dtype: int64
```

## Perform Correlation Analysis

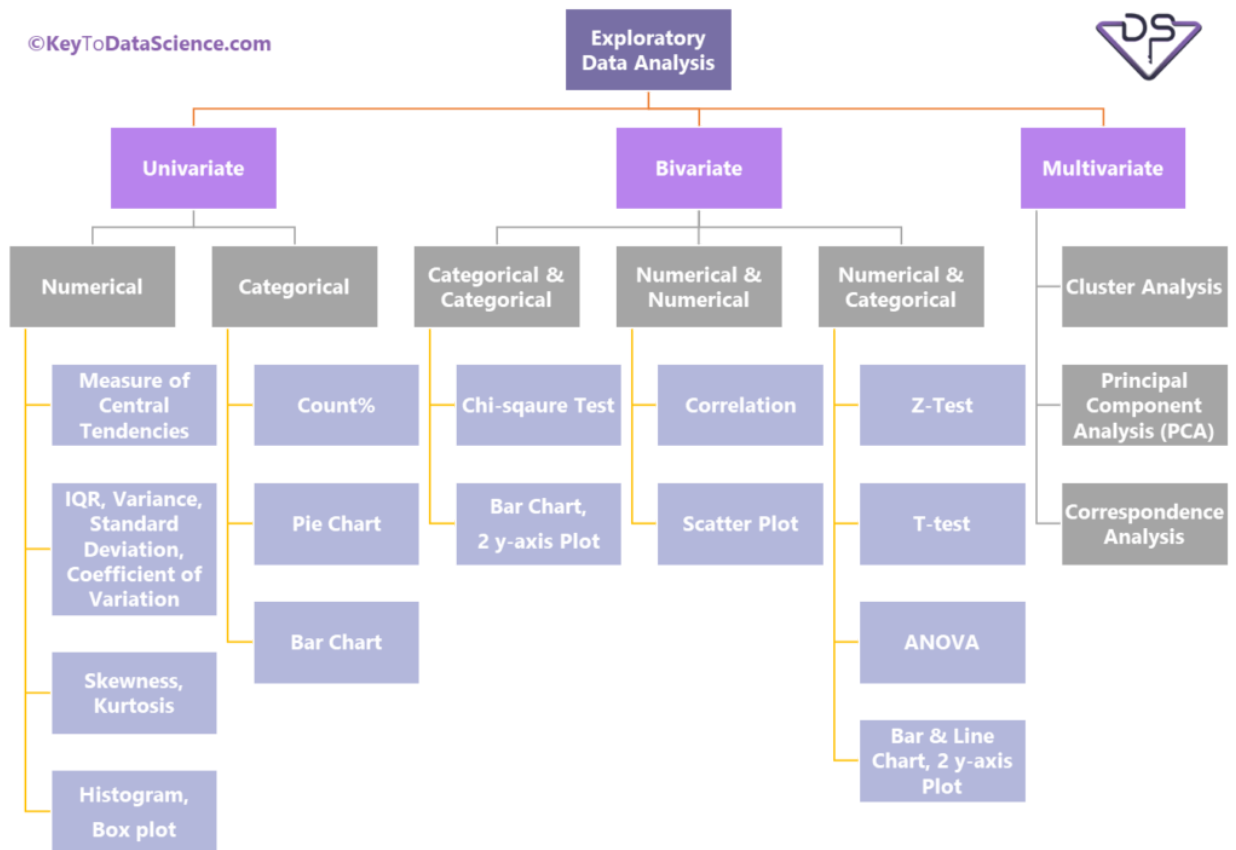
```
In [9]: df.corr()['CLV']
        # -1 to 1
```

```
Out[9]: MONTH_1    0.734122
        MONTH_2    0.250397
        MONTH_3    0.371742
        MONTH_4    0.297408
        MONTH_5    0.376775
        MONTH_6    0.327064
        CLV        1.000000
        Name: CLV, dtype: float64
```

We can see that the months do show strong correlation to the target variable (CLV). That should give us confidence that we can build a strong model to predict the CLV

## 3.2 Exploratory Data Analysis (EDA)

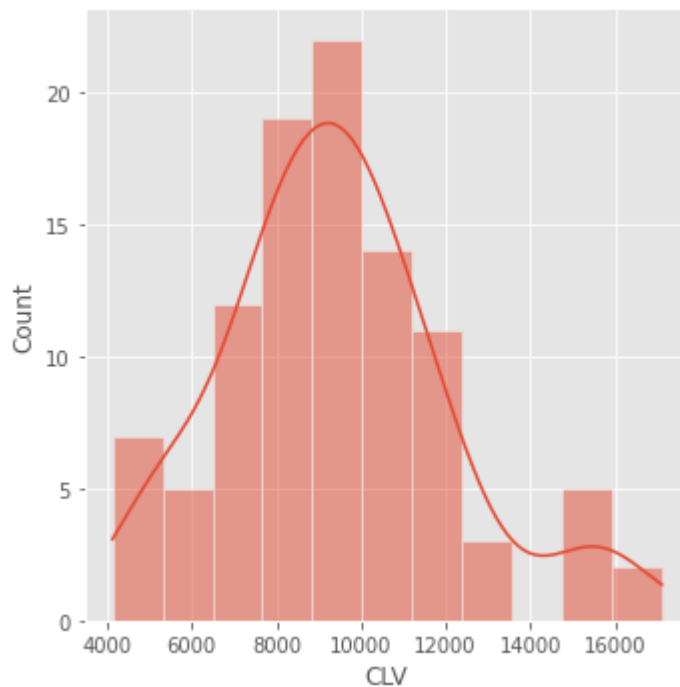
### Exploratory Data Analysis (EDA)<sup>3</sup>



## Univariate Analysis <sup>4</sup>

```
In [10]: sns.displot(df['CLV'],kde=True)
```

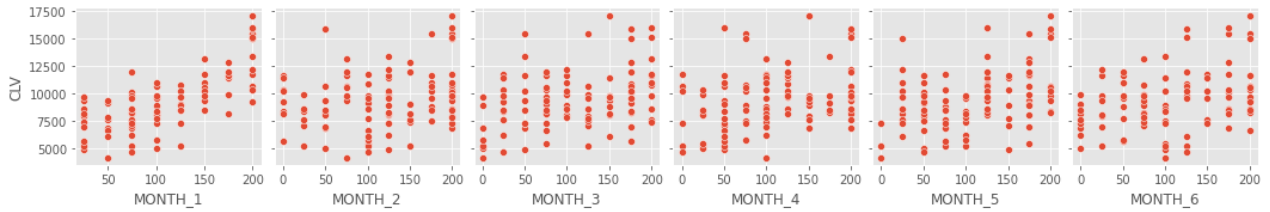
```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x18521bcfb50>
```



## Bivariate Analysis <sup>5</sup>

```
In [11]: # df.columns[:-1]
sns.pairplot(df,x_vars=df.columns[:-1],y_vars=['CLV'])
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x18525cf9fd0>
```



## 3.3 Train and Test Split

Prepare X(independent) and y(dependent) variables

```
In [12]: #Drop columns with low correlation
X = df.drop('CLV',axis=1)
y = df.CLV
```

Let us split the data into training and testing datasets in the **ratio 70:30**

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3)
print("X_train - Training : ", X_train.shape)
print("X_test - Testing : ", X_test.shape)
```

```
X_train - Training : (70, 6)
X_test - Testing : (30, 6)
```

[↑ back to top](#)

## 4. Modeling

Create a function to quickly build models and evaluate

```
In [14]: def model_builder(algo,X_train,y_train,X_test):
    algo.fit(X_train,y_train)

    y_train_pred = algo.predict(X_train)
    y_test_pred = algo.predict(X_test)

    print("The model performance for training set")
    print("-----")
    print(f'R2 : {round(r2_score(y_true=y_train,y_pred=y_train_pred),2)}')
    print('MAE :', mean_absolute_error(y_train, y_train_pred))
    print('MSE :', mean_squared_error(y_train, y_train_pred))
    print('RMSE:', np.sqrt(mean_squared_error(y_train, y_train_pred)))

    print("\n")

    print("The model performance for testing set")
    print("-----")
    print(f'R2 : {round(r2_score(y_true=y_test,y_pred=y_test_pred),2)}')
    print('MAE :', mean_absolute_error(y_test, y_test_pred))
    print('MSE :', mean_squared_error(y_test, y_test_pred))
    print('RMSE:', np.sqrt(mean_squared_error(y_test, y_test_pred)))
```

## 4.1 Linear Regression <sup>6</sup>

### Build Model

We build a Linear Regression equation for predicting CLV and then check its accuracy by predicting against the test dataset

```
In [15]: # #Build model on training data
# lr = LinearRegression()
# lr.fit(X_train,y_train)
# print("Coefficients :", lr.coef_)
# print("Intercept    :", lr.intercept_)
```

```
In [16]: # use model_builder function
lr = LinearRegression()
model_builder(lr,X_train,y_train,X_test)
```

The model performance for training set

-----

R2 : 0.93  
MAE : 595.5826744945408  
MSE : 509998.45508921274  
RMSE: 714.1417611995623

The model performance for testing set

-----

R2 : 0.9  
MAE : 637.6508538506049  
MSE : 598689.529200415  
RMSE: 773.7503015834081

## 4.2 DecisionTreeRegressor <sup>7</sup>

```
In [17]: # from sklearn.tree import DecisionTreeRegressor

dec_tree = DecisionTreeRegressor(random_state=0)
model_builder(dec_tree,X_train,y_train,X_test)
```

The model performance for training set

-----

R2 : 1.0  
MAE : 0.0  
MSE : 0.0  
RMSE: 0.0

The model performance for testing set

-----

R2 : 0.57  
MAE : 1286.8333333333333  
MSE : 2599523.3666666667  
RMSE: 1612.3037451630096



## 4.3 RandomForestRegressor <sup>8</sup>

```
In [18]: # from sklearn.ensemble import RandomForestRegressor

rf_tree = RandomForestRegressor(random_state=0)
model_builder(rf_tree,X_train,y_train,X_test)
```

The model performance for training set

-----

R2 : 0.97  
MAE : 392.26885714285714  
MSE : 260275.18048571429  
RMSE: 510.17171666578525

The model performance for testing set

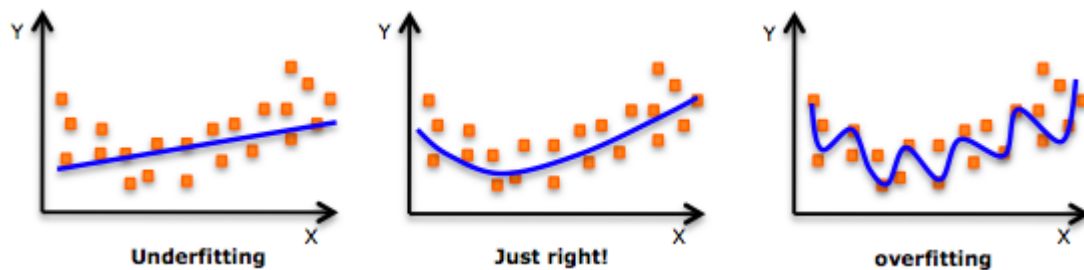
-----

R2 : 0.75  
MAE : 920.8510000000002  
MSE : 1483953.4223700003  
RMSE: 1218.1762690062553

### Observation:

Linear Regression model has the highest: R2; and lowest: MAE, MSE, RMSE on the test set.

### Both Decision Tree Regressor and Random Forest Regressor are overfitting



Reference Image for Overfitting

[↑ back to top](#)

## 5. Evaluation <sup>9</sup>

### 📌 Regression Evaluation Metrics

Here are three common evaluation metrics for regression problems:

**Mean Absolute Error (MAE)** is the mean of the absolute value of the errors:

$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

**Mean Squared Error (MSE)** is the mean of the squared errors:

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

**Root Mean Squared Error** (RMSE) is the square root of the mean of the squared errors:

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Comparing these metrics:

- **MAE** is the easiest to understand, because it's the average error.
- **MSE** is more popular than MAE, because MSE "punishes" larger errors, which tends to be useful in the real world.
- **RMSE** is even more popular than MSE, because RMSE is interpretable in the "y" units.

All of these are **loss functions**, because we want to minimize them.

## 5.1 Select Final Model

**Final Model Selected for production is:** Linear Regression

Linear Regression model has the highest: R2; and lowest: MAE, MSE, RMSE on the test set.

In [19]:

```
y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)

print("The model performance for training set")
print("-----")
print(f'R2 : {round(r2_score(y_true=y_train,y_pred=y_train_pred),2)}')
print('MAE :', mean_absolute_error(y_train, y_train_pred))
print('MSE :', mean_squared_error(y_train, y_train_pred))
print('RMSE:', np.sqrt(mean_squared_error(y_train, y_train_pred)))

print("\n")

print("The model performance for testing set")
print("-----")
print(f'R2 : {round(r2_score(y_true=y_test,y_pred=y_test_pred),2)}')
print('MAE :', mean_absolute_error(y_test, y_test_pred))
print('MSE :', mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(mean_squared_error(y_test, y_test_pred)))
```

The model performance for training set

```
-----
R2 : 0.93
MAE : 595.5826744945408
MSE : 509998.45508921274
RMSE: 714.1417611995623
```

The model performance for testing set

```
-----
R2 : 0.9
MAE : 637.6508538506049
```

MSE : 598689.529200415  
RMSE: 773.7503015834081

Linear Regression shows a R-squared of 0.91% on Testing set. This is an excellent model for predicting CLV.

## 5.2 Save Model to Disk

```
In [20]: # get current directory path  
# os.getcwd()
```

```
In [21]: # import joblib  
filename = os.getcwd()+'/CLV_LinearRegression.joblib.pkl'  
joblib.dump(lr, filename, compress=9)
```

```
Out[21]: ['F:\\Work\\Site\\KDS - Career Now Program\\DS\\Syllabus\\5. Case Studies\\Business\\Customer Lifetime Value (CLV)/CLV_LinearRegression.joblib.pkl']
```

## 5.3 Interpret the Output

```
In [22]: # print the intercept  
print("Intercept:", lr.intercept_)
```

Intercept: -250.46672566544476

```
In [23]: coeff_df = pd.DataFrame(lr.coef_, X.columns, columns=['Coefficient'])  
coeff_df
```

```
Out[23]:
```

	Coefficient
<b>MONTH_1</b>	34.765858
<b>MONTH_2</b>	10.694871
<b>MONTH_3</b>	15.682293
<b>MONTH_4</b>	12.583314
<b>MONTH_5</b>	8.425176
<b>MONTH_6</b>	5.615516

### Interpreting the coefficients:

- Holding all other features fixed, a 1 unit increase in **MONTH\_1** is associated with an **increase of \$34.7** in Customer Lifetime Value (CLTV).
- Holding all other features fixed, a 1 unit increase in **MONTH\_2** is associated with an **increase of \$10.6** in Customer Lifetime Value (CLTV).
- Holding all other features fixed, a 1 unit increase in **MONTH\_3** is associated with an **increase of \$15.6** in Customer Lifetime Value (CLTV).
- Holding all other features fixed, a 1 unit increase in **MONTH\_4** is associated with an **increase of \$12.5** in Customer Lifetime Value (CLTV).

- Holding all other features fixed, a 1 unit increase in **MONTH\_5** is associated with an **increase of \$8.4** in Customer Lifetime Value (CLTV).
- Holding all other features fixed, a 1 unit increase in **MONTH\_6** is associated with an **increase of \$5.6** in Customer Lifetime Value (CLTV).

## 5.4 Linear Regression with StandardScaler (Optional)

```
In [24]: # from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaling = scaler.fit_transform(X_train)
```

💡 To preventing information about the distribution of the test set leaking into the model.

**Fit the scaler on your training data only, then standardise both training and test sets with that scaler.**

```
In [25]: # we have to scale X_test also, before predicting
# Use transform() on the test data, not fit_transform(), as fit is done on training set
X_test_scaling = scaler.transform(X_test)
```

```
In [26]: lr_scaling = LinearRegression()
model_builder(lr_scaling,X_train_scaling,y_train,X_test_scaling)
```

The model performance for training set

```
-----
R2   : 0.93
MAE  : 595.5826744945401
MSE  : 509998.4550892126
RMSE : 714.1417611995623
```

The model performance for testing set

```
-----
R2   : 0.9
MAE  : 637.6508538506054
MSE  : 598689.5292004154
RMSE : 773.7503015834084
```

```
In [27]: coeff_df_scaling = pd.DataFrame(lr_scaling.coef_,X.columns,columns=['Coefficient'])
coeff_df_scaling
```

```
Out[27]:
```

	Coefficient
<b>MONTH_1</b>	1964.026074
<b>MONTH_2</b>	678.352205
<b>MONTH_3</b>	980.639262
<b>MONTH_4</b>	773.654048
<b>MONTH_5</b>	501.735199

	Coefficient
MONTH_6	368.845216

[↑ back to top](#)

## 6. Model Deployment

The code in model deployment should be able to run independently, i.e. the below code should be independent of all the above performed steps.

### Predicting for a new Customer

Let's use the model to predict his CLV.

### 6.1 Import Libraries

```
In [28]: # import os
# import joblib
# from sklearn.linear_model import LinearRegression
```

### 6.2 Load Model from Disk

```
In [29]: filename = os.getcwd()+'/CLV_LinearRegression.joblib.pkl'
```

```
In [30]: model = joblib.load(filename)
```

### 6.3 Real Time Prediction

#### New Customer Data

Say we have a new customer who in his first 3 months have spend **100,0,50** on the website.

```
In [31]: new_data = np.array([100,0,50,0,0,0]).reshape(1, -1)
```

#### Real Time Prediction

```
In [32]: new_pred=model.predict(new_data)
print("The CLV for the new customer is : $",new_pred[0])
```

The CLV for the new customer is : \$ 4010.2337033822596

[↑ back to top](#)

Great Job!

---

Links Used in this notebook:

1. [Data Science Project Life Cycle](#) <sup>1</sup>
2. [Customer Lifetime Value\(CLTV\)](#) <sup>2</sup>
3. [Exploratory Data Analysis \(EDA\)](#) <sup>3</sup>
4. [Univariate Analysis](#) <sup>4</sup>
5. [Bivariate Analysis](#) <sup>5</sup>
6. [Regression](#) <sup>6</sup>
7. [Decision Tree](#) <sup>7</sup>
8. [Random Forest](#) <sup>8</sup>
9. [Model Evaluation](#) <sup>9</sup>

**[KeytoDataScience.com](https://keytoDataScience.com)**